

# Forcing Correspondence (v16): Calibrated Structural Transduction Between Set-Theoretic Forcing and the Hypothesis Surface

Paul Tiffany

July 21, 2026

## Abstract

We give a structural correspondence between set-theoretic (Cohen) forcing and the refinement dynamics of a resource-bounded hypothesis surface. Geometric forcing is *not* identical to Cohen forcing; the claim is a conditional transduction theorem. Both instantiate a common *forcing-over-a-site* schema: the forcing relation is *dense stabilization*—a condition forces a proposition when the refinements stabilizing it are dense below it in the channel-margin subposet—and the operationally admissible dense requirements *generate* a Grothendieck topology  $J_{\text{adm}}$  on the channel-margin subposet, sandwiched below the double-negation topology, so the Kripke–Joyal interpretation transfers clause by clause. The site is the closure of obligation under refinement: raw admissible obligations  $\mathcal{D}_{\text{adm}}$  generate the semantic site  $J_{\text{adm}}$ , of which a finite-prefix *witnessed* fragment  $J_{\text{wit}} \subseteq J_{\text{adm}}$  is the executable discipline. Channel-margin preservation is a path-budget perturbation lemma (the per-step form does not iterate); separativity is carried by certificate disagreement, not by the metric margin; generic existence is Rasiowa–Sikorski under the observer bounded-countability condition; the generic is read geometrically as a Cauchy-generic completion in an observer completion, with an order–metric compatibility lemma tying the generic filter to the limit state, conditional on a smooth descent contract coextensive with the channel margin. Atomic soundness follows from eventual stabilization; propositional soundness follows by the standard forcing induction, conditional on a decision-reachability lemma and a calibration hypothesis here decomposed into a finite theorem queue rather than asserted. Inter-observer transfer is restored as an essential interface medium  $M$ : the void supplies commensurability; a public medium and protocol transport only encoded residues  $E_A(T_A x_A)$ , and reconstruction is generically not identity—an equality-condition theorem. The executable interpretation is witness-only PyLantern/Chalked traces. Quantifiers, fullness/witness clauses, and implementation calibration remain open; every load-bearing claim is tagged P/D/S/M/C/O in a status ledger (§16).

## 1 Two senses of “forcing”

The word is currently overloaded across this program; separating the two uses is notation, not a retraction.

*Control forcing* (Hypothesis Surface §4.1) is the precondition: the geometric cost of joint progress on  $k$  claims diverges at the conflict threshold, so a bounded observer is forced, dynamically, to stop differentiating and integrate. This is the event that deposits the observer onto a refinement poset.

*Logical forcing* ( $\Vdash_H$ , this note) is the relation that then lives on that poset. We reserve “forces” for the second and “forced transition” for the first.

## 2 Core definitions (instantiated)

**Definition 1** (Conditions).  $P_H$  is the set of claim complexes  $K_t$  with fibers  $f_i = (\sigma_i, r_i, W_i)$  and exit cone  $E(x_0)$ . A condition is a finite partial epistemic state.

**Assumption 1** (Smoothness contract). Fix a margin  $\eta > 0$  and an anchor condition  $p_0$  with  $\lambda_{\min}(\tilde{\Gamma}(p_0)) \geq \eta$ . Along admissible trajectories from  $p_0$ : (a)  $\tilde{\Gamma}$  is  $L_\Gamma$ -Lipschitz in spectral norm with respect to the projected state displacement; (b) the cumulative displacement obeys the budget  $\sum_n \|\Delta_n\| \leq \eta/(2L_\Gamma)$ . This is the displacement contract of Cacophony §2.1 in path form; it is *measured, not assumed* (see the remark following the Channel-Margin Lemma in §7), and it is the single calibration debt written “(ii)” in earlier versions of this note.

**Definition 2** (Refinement order and channel-margin subposet).  $q \leq p$  iff  $q$  is reachable from  $p$  by an admissible move retaining  $p$ ’s commitments. The *channel-margin subposet*  $\mathbb{P}_H^\eta$  consists of the conditions reachable from the anchor  $p_0$  by admissible moves within the displacement budget of Assumption 1. All forcing below is relative to  $\mathbb{P}_H^\eta$ . (That  $\lambda_{\min}(\tilde{\Gamma}) \geq \eta/2$  holds throughout  $\mathbb{P}_H^\eta$  is proved, not stipulated, by the Channel-Margin Lemma of §7; the definition consumes only the budget.)

**Definition 3** (Sieves and dense requirements). For  $p \in \mathbb{P}_H^\eta$  a *sieve* on  $p$  is a downward-closed subset  $S \subseteq \downarrow p = \{q \in \mathbb{P}_H^\eta : q \leq p\}$ . A family  $D$  of conditions below  $p$  is an *admissible dense requirement* if for every  $q \leq p$  in  $\mathbb{P}_H^\eta$  there is a further *admissible* refinement  $r \leq q$  in  $\mathbb{P}_H^\eta$  meeting  $D$  (the paradigm cases are certificate transitions and voids, but the class is not restricted to them—deciding families must be eligible); the sieve it generates is  $S_D = \{q \leq p : \exists \text{ admissible } r \leq q, r \in D\}$ . Only admissible dense requirements count; not every void or certificate the observer names qualifies. Write  $\mathcal{D}_{\text{adm}}$  for the class of all admissible dense requirements on  $\mathbb{P}_H^\eta$ .

**Definition 4** (Stability, atomic). For *atomic*  $\varphi$  in the enrolled vocabulary,  $\text{Stab}(q, \varphi)$  holds when  $\varphi$  holds at  $q$  and persists under further admissible refinement (the soft-min propagation of Eq. 1 makes confidence monotone down the order; that Eq. 1 propagates this way is a modeling postulate about the apparatus—the highest-leverage postulate in this note, consumed by *every* truth lemma below, tracked as (M-Pers) in §16). The *stability set* of atomic  $\varphi$  below  $p$  is

$$A_\varphi(p) = \{q \leq p : \text{Stab}(q, \varphi)\}.$$

By (M-Pers),  $A_\varphi(p)$  is closed under further refinement, hence a sieve on  $p$ .  $\text{Stab}$  is *not* defined for compound formulas; compound forcing is defined recursively by the propositional clauses of §5, and the torsion remark of §7 records where a compound- $\text{Stab}$  reading would diverge from them.

**Definition 5** (Preservation set).  $\mathcal{I} = \{\text{I1 No Bare Claims (well-definedness), I2 Anti-Masking (boundedness: keeps } \tilde{\Gamma} \text{ faithful, hence } \lambda_{\min} \text{ honest), I3 Certificate Classification (obstruction typing)}\}$ .

## 3 The generated admissible site: raw, generated, witnessed

The forcing-over-refinement schema needs a site. We do *not* claim that the raw operational class  $\mathcal{D}_{\text{adm}}$  is already closed under the Grothendieck operations—raw admissibility depends on channel margin, smoothness, finite witnessability, resource bounds, and protocol constraints, and may fail pullback/transitivity closure as an operational class. Instead we *generate* the topology.

**Definition 6** (Generated admissible site). Let  $B = \{S_D : D \in \mathcal{D}_{\text{adm}}\}$  be the collection of sieves generated by admissible dense requirements (Def. 3). Let

$$J_{\text{adm}} = \langle B \rangle$$

be the least Grothendieck topology on the poset  $\mathbb{P}_H^\eta$  containing  $B$ . The pair  $(\mathbb{P}_H^\eta, J_{\text{adm}})$  is the *generated admissible site*.

**Lemma 1** (Site bound).  $J_{\text{adm}}$  exists and is a Grothendieck topology, and  $J_{\text{adm}} \subseteq J_{\neg\neg}$ , where  $J_{\neg\neg}$  is the double-negation (dense) topology whose covers are exactly the dense sieves. In particular  $(\mathbb{P}_H^\eta, J_{\text{adm}})$  is a genuine site.

*Proof.* The least Grothendieck topology containing an arbitrary collection of sieves always exists: the intersection of all Grothendieck topologies containing  $B$  is again a Grothendieck topology and is the least such. For the bound, recall that on any poset the dense sieves form a Grothendieck topology, the double-negation topology  $J_{\neg\neg}$  (a sieve  $S$  on  $p$  is  $J_{\neg\neg}$ -covering iff for every  $q \leq p$  some  $r \leq q$  lies in  $S$ ; see [4]). Each generator  $S_D \in B$  is dense below  $p$  by admissible density (Def. 3), hence  $S_D \in J_{\neg\neg}$ . Since  $J_{\neg\neg}$  is a Grothendieck topology containing  $B$ , the least such is contained in it:  $J_{\text{adm}} \subseteq J_{\neg\neg}$ .  $\square$

**Definition 7** (Forcing relations). For  $p \in \mathbb{P}_H^\eta$  and atomic  $\varphi$  (stability sets  $A_\varphi(p)$ , Def. 4):

$$p \Vdash_H \varphi :\iff A_\varphi(p) \in J_{\text{adm}}(p), \quad p \Vdash_{\neg\neg} \varphi :\iff A_\varphi(p) \text{ is dense below } p \text{ (a } J_{\neg\neg}\text{-cover)}.$$

$\Vdash_H$  is the *calibrated* relation of this note;  $\Vdash_{\neg\neg}$  is the ambient dense-below reading, the characterization of Cohen forcing transported onto  $K_t$ . By Lemma 1,  $p \Vdash_H \varphi \Rightarrow p \Vdash_{\neg\neg} \varphi$ . The two relations are distinguished notationally throughout; each lemma below states which relation it is about, and transfer between them is never silent (it is exactly the Decision Reachability question). Both extend to compound formulas by the recursive clauses of §5. The strictly stronger condition “ $\text{Stab}(q, \varphi)$  for every  $q \leq p$ ” (the maximal sieve, a cover in every topology) forces under both relations and is reserved for torsion (§7, Remark).

*Remark* (The site is the closure of obligation under refinement). Rigor does not flatten admissibility; it completes it. The operational layer supplies admissible dense *obligations*  $\mathcal{D}_{\text{adm}}$ ; the mathematical layer closes them under the Grothendieck operations into the site  $(\mathbb{P}_H^\eta, J_{\text{adm}})$ . Kripke–Joyal forcing over  $(\mathbb{P}_H^\eta, J_{\text{adm}})$  supplies the abstract forcing semantics (§5).  $\Vdash_{\neg\neg}$  is the Kripke–Joyal reading for the ambient  $J_{\neg\neg}$ ; the calibrated relation  $\Vdash_H$  uses the generated subtopology  $J_{\text{adm}} \subseteq J_{\neg\neg}$ , hence is *at least as strong*. Whether a given deciding sieve is actually a  $J_{\text{adm}}$ -cover is the Decision Reachability question (Lemma 7).

**Definition 8** (Witnessed fragment). The *witnessed fragment*  $J_{\text{wit}} \subseteq J_{\text{adm}}$  is the sub-collection of  $J_{\text{adm}}$ -covers that admit a finite admissible protocol witness: a finite prefix of refinements, each carrying a certificate, admissibility check, provenance, and (where required) a human-gated decision, whose density below  $p$  is verified by that finite data.  $J_{\text{wit}}$  is in general *not* a topology; it is the executable basis. The three layers

$$\mathcal{D}_{\text{adm}} \rightsquigarrow J_{\text{adm}} \rightsquigarrow J_{\text{wit}} \subseteq J_{\text{adm}}$$

read: raw operational obligations  $\rightsquigarrow$  generated semantic closure  $\rightsquigarrow$  finite-prefix, PyLantern/Chalked-witnessable fragment.

*Remark* (Self / other / global). The three layers are the self/other/global stack. Raw is what the self emits; generated is what the other may require; witnessed is what the global surface can audit. A hypothesis leaves the self as a raw obligation, meets the other as generated refinement pressure, and stabilizes globally only as witnessed residue. Formal closure must not silently authorize operational traversal: *the generated site supplies semantics; the witnessed fragment supplies executable discipline* ( $J_{\text{adm}} \supseteq J_{\text{wit}}$  is a containment, not an equality, and §13 operates only in  $J_{\text{wit}}$ ). This is exactly why Chalked (§14) matters: it is the finite witness surface on which  $J_{\text{wit}}$  becomes consent-gated and inspectable.

## 4 Correspondence table

| Component                      | Cohen / set-theoretic                      | Hypothesis Surface (concrete)                                                                     |
|--------------------------------|--------------------------------------------|---------------------------------------------------------------------------------------------------|
| Condition                      | Finite partial function                    | Claim complex $K_t + \text{fibers}$ ; exit cone $E(x_0)$                                          |
| Refinement $q \leq p$          | Stronger condition                         | Admissible move in $\mathbb{P}_H^\eta$                                                            |
| Channel viability              | (nondegeneracy of extensions)              | $\lambda_{\min}(\tilde{\Gamma}) > 0$ (off-cliff; $\rho(k-1) < 1$ )                                |
| Separativity / incompatibility | Incompatible extensions below $p$          | <i>Certificate disagreement</i> (I3); undecided obstruction (Lem 5)                               |
| Requirement / cover            | Dense set of conditions                    | Admissible dense requirement (Def. 3); $J_{\text{adm}}$ -cover (Def. 6)                           |
| Generic / completion           | Generic filter $G$                         | Admissible-complete trajectory in $\mathbb{P}_H^\eta$                                             |
| Forcing $p \Vdash \varphi$     | Forcing-set dense below $p$                | $A_\varphi(p)$ a $J_{\text{adm}}$ -cover of $p$ ( $\Vdash_H$ ); dense form is $\Vdash_{\neg\neg}$ |
| Extension                      | $M[G]$                                     | Stabilized surface along the completed trajectory                                                 |
| Preservation                   | Ground theory preserved                    | Invariants I1–I3 survive completion                                                               |
| Forced negation                | $p \Vdash \neg\varphi$                     | Torsion certificate (atomic $\varphi$ refuted on <i>all</i> refinements; §7)                      |
| Undecided                      | $p \nVdash \varphi, p \nVdash \neg\varphi$ | Non-torsion certificate (contingent)                                                              |
| Soundness                      | Truth Lemma                                | Atomic + Propositional Truth Lemma (Thm. 1); quantifiers open                                     |
| Shared zero / origin           | Top condition $\nVdash$                    | Void (Axiom H.0); unique shared depth-0 condition                                                 |

**Separativity correction.** The metric margin and separativity are distinct roles: *margin keeps the channel open; certificates make branches incompatible*. The convex metric part of every exit cone is self-compatible (Lemma 5, using convexity), so  $\lambda_{\min} > 0$  supplies channel *viability*—the existence of admissible refinements—not separativity. Incompatibility, hence the separativity needed for a nondegenerate forcing poset, is carried by the discrete certificate axis (I3): two refinements are incompatible exactly when they disagree on an undecided obstruction.

## 5 Formal schema

**Forcing schema.**  $p \Vdash_H \varphi$  when the refinements of  $p$  that stabilize  $\varphi$  form a  $J_{\text{adm}}$ -cover of  $p$  in the channel-margin subposet.

*Remark* (Rigorous home). The schema is the persistence-monotone (intuitionistic) reading of forcing on a site: Kripke–Joyal forcing semantics over  $(\mathbb{P}_H^\eta, J_{\text{adm}})$ , with admissible dense requirements generating the covering families. Cohen forcing is the set-theoretic instance; its topos presentation routes through Kripke–Joyal semantics and the double-negation/classicalization construction [4], i.e. through  $J_{\neg\neg}$ , of which  $J_{\text{adm}}$  is the generated subtopology (Lemma 1). “Calibrated structural transduction” is then the claim that both Cohen forcing and the Hypothesis Surface instantiate forcing-over-a-site, the calibration being which requirements generate the site, and the preserved invariant being the Kripke–Joyal interpretation of the connectives. The transfer is verifiable clause by clause.

**Definition 9** (Propositional clauses and deciding sets). Let  $J \in \{J_{\text{adm}}, J_{\neg\neg}\}$  and write  $\Vdash$  for the corresponding atomic relation of Def. 7. Take  $\neg, \wedge, \vee$  primitive over  $\mathbb{P}_H^\eta$  and define  $\varphi \rightarrow \psi := \neg\varphi \vee \psi$ :

$$\begin{aligned} p \Vdash \varphi \wedge \psi &\iff p \Vdash \varphi \text{ and } p \Vdash \psi; \\ p \Vdash \varphi \vee \psi &\iff \{q \leq p : q \Vdash \varphi \text{ or } q \Vdash \psi\} \in J(p); \\ p \Vdash \neg\varphi &\iff \text{no } q \leq p \text{ has } q \Vdash \varphi. \end{aligned}$$

The last is the universal-refinement clause; at the void it is the across-every-branch statement of §9. The *deciding set* of  $\varphi$  is  $D_\varphi = \{q : q \Vdash \varphi \text{ or } q \Vdash \neg\varphi\}$  (relation as per context;  $\Vdash_H$  unless marked).

**Lemma 2** (Persistence and consistency). *For  $\Vdash \in \{\Vdash_H, \Vdash_{\neg\neg}\}$  and every propositional  $\varphi$ : if  $p \Vdash \varphi$  and  $q \leq p$  then  $q \Vdash \varphi$ ; and no  $p$  forces both  $\varphi$  and  $\neg\varphi$ .*

*Proof.* Persistence by induction on  $\varphi$ : atomic because  $A_\varphi(q) = A_\varphi(p) \cap \downarrow q$  is the pullback of the sieve  $A_\varphi(p)$  along  $q \leq p$ , and covers pull back in any Grothendieck topology (for  $J_{\neg\neg}$ : dense restricts to dense);  $\wedge$  componentwise;  $\vee$  because a  $J$ -cover of  $p$  restricts to a  $J$ -cover of any  $q \leq p$  (stability axiom);  $\neg$  because its defining cone is downward-closed. Consistency:  $p \Vdash \neg\varphi$  says no  $q \leq p$  forces  $\varphi$ , yet  $p \leq p$ , so  $p \not\Vdash \varphi$ . Beyond the well-formedness of the atomic relation ( $A_\varphi$  a sieve, which is (M-Pers), Def. 4), clause-level persistence consumes no additional postulate—it is pullback stability of the topology; (M-Pers) does further substantive work only in the Atomic Truth Lemma.  $\square$

**Lemma 3** (Deciding sets are order-dense). *For  $\Vdash \in \{\Vdash_H, \Vdash_{\neg\neg}\}$  and every propositional  $\varphi$ , the deciding set  $D_\varphi$  is dense in the order.*

*Proof.* Given  $p$ : if some  $q \leq p$  forces  $\varphi$  then  $q \in D_\varphi$ ; otherwise no  $q \leq p$  forces  $\varphi$ , which is exactly  $p \Vdash \neg\varphi$ , so  $p \in D_\varphi$ . Density is the  $\neg$ -clause read contrapositively—uniform in  $\varphi$  and in the choice of relation.  $\square$

*Remark.* Order-density of  $D_\varphi$  is free; whether  $D_\varphi$  is *admissibly* dense—a  $J_{\text{adm}}$ -cover—is not, and is precisely the Decision Reachability question of §7. By Lemma 2,  $D_\varphi$  is closed under refinement, hence a sieve on any  $p$  when restricted below it.

## 6 Calibration as a finite theorem queue

The forcing lemmas below are theorems about the abstract site semantics. Their bearing on the Hypothesis Surface pipeline requires a calibration hypothesis: the concrete surface operations must realize the Kripke–Joyal clauses they are claimed to instantiate. Calibration is not one foggy assumption; it is a *finite theorem queue*, each item with a definite home.

| #  | Calibration lemma            | Home / status                                                                                                                                         |
|----|------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1  | Atomic Stabilization         | claim stabilization = Stab; <i>Hypothesis Surface</i>                                                                                                 |
| 2  | Refinement Audit             | admissible move = refinement, audit-visible; <i>PyLantern</i>                                                                                         |
| 3  | Dense Requirement            | operational requirements admissibly dense; measurable (C, Assumption 1)                                                                               |
| 4  | Certificate Conjunction      | certificate combination = $\wedge$ -clause; <i>Hypothesis Surface</i>                                                                                 |
| 5  | Admissible-Cover Disjunction | branching = dense-cover $\vee$ -clause; <i>Hypothesis Surface</i>                                                                                     |
| 6  | Torsion Negation             | torsion certificate = universal-refinement $\neg$ ; discharged for <i>atomic</i> $\varphi$ (§8, Remark); compound case <b>O</b> (finite countermodel) |
| 7  | Soft-Min Persistence         | Eq. 1 = monotone persistence; <i>Hypothesis Surface</i> (M-Pers)                                                                                      |
| 8  | Decision Reachability        | deciding sieve is a $J_{\text{adm}}$ -cover; Lem 7, conditional on Assumption 1                                                                       |
| 9  | Observer Completion          | $(P_O, d_O)$ metric, completion exists; <i>Principia</i>                                                                                              |
| 10 | Interface Reconstruction     | $E_A, D_B$ carry witness/provenance/consent, $D_B \circ E_A \neq \text{id}$ ; <i>PyLantern/Chalked</i> + Thm 3                                        |

Items 1, 4, 5, 7 are supplied by *The Hypothesis Surface*; 2, 10 by PyLantern/Chalked implementation; 3, 8 reduce to the measurable smoothness contract (Assumption 1); 6 is discharged here for atomic formulas only; 9 is an interface to *Principia*. *None is a fresh unknown*. The Propositional Truth Lemma is operationally meaningful under this queue; discharging it is more urgent than adding quantifiers, because it is the check that turns structural analogy into verified transduction for the actual pipeline.

**Assumption 2** (Calibration queue). The realization lemmas 1–10 of the table above hold for the deployed pipeline: the concrete surface operations realize the Kripke–Joyal clauses they are claimed to instantiate. Items already discharged or reduced retain their homes as listed; the assumption is the conjunction of the remainder. Status **O** overall; consumed explicitly wherever a result is claimed *about the pipeline* rather than about the site semantics.

## 7 Theorem target, with proofs of the joints

**Theorem Target.** *Let  $H$  be a hypothesis surface carried by a resource-bounded observer, with claim complex  $K_t$ , interaction matrix  $\tilde{\Gamma}_t = G_t G_t^\top$ , and channel projector  $T_t$ . Suppose: (1)  $\lambda_{\min}(\tilde{\Gamma}_t(p)) \geq \eta > 0$  at the anchor; (2) admissible trajectories satisfy the smoothness contract (Assumption 1), so the Channel-Margin Lemma (Lemma 4) applies; (3)  $\mathcal{D}_p$  is a countable family of admissible dense requirements generating (with the deciding closure)  $J_{\text{adm}}$ -covers, with decision reachability (Lemma 7); (4) the observer-completion axioms supply Cauchy convergence for admissible refinement towers (Lemma 8) and order-metric compatibility (Lemma 9). Then there is a Cauchy-generic trajectory  $G$  through  $p$  meeting every requirement and converging to  $x_G \in \bar{P}_O$ ; along  $G$ ,  $\Vdash_H$  is monotone,  $\mathcal{I}$  is preserved, atomic forced propositions are exactly those eventually stable (Lemma 10), and soundness extends to the propositional connectives (Thm. 1) under the calibration queue.*

1. **Monotonicity (persistence).**  $p \Vdash_H \varphi \Rightarrow q \Vdash_H \varphi$  for  $q \leq p$ , by pullback stability of covers (Lemma 2); no apparatus postulate is consumed.
2. **Channel viability, persistence, branching (Lemmas 4, 5).**  $\lambda_{\min}(\tilde{\Gamma}) > 0$  is the metric non-degeneracy making admissible refinements exist (viability); the Channel-Margin Lemma keeps the margin positive; separativity/branching is carried by certificate disagreement (Lemma 5), not the convex metric.
3. **Generic existence (Lemma 6 + Rasiowa–Sikorski).**
4. **Decision reachability (Lemma 7).**
5. **Atomic soundness (Lemma 10); propositional soundness (Thm. 1).** Open layers: quantifiers, and the calibration queue (§6).

**Lemma 4** (Channel-Margin, path form). *Let  $\lambda(q) = \lambda_{\min}(\tilde{\Gamma}(q))$  and let  $q$  be reachable from the anchor  $p_0$  by an admissible trajectory with steps  $\Delta_1, \dots, \Delta_n$ . Under Assumption 1 ( $\lambda(p_0) \geq \eta$ ;  $L_\Gamma$ -Lipschitz; budget  $\sum_i \|\Delta_i\| \leq \eta/(2L_\Gamma)$ ),*

$$\lambda(q) \geq \lambda(p_0) - L_\Gamma \sum_{i \leq n} \|\Delta_i\| \geq \eta/2 > 0.$$

Hence  $\mathbb{P}_H^\eta$  (Def. 2) stays in the channel regime, uniformly in the depth of the descent.

*Proof.*  $\lambda_{\min}$  is 1-Lipschitz in the spectral norm (Weyl):  $|\lambda(q) - \lambda(p_0)| \leq \|\tilde{\Gamma}(q) - \tilde{\Gamma}(p_0)\|$ . By the Lipschitz hypothesis and the triangle inequality along the trajectory,  $\|\tilde{\Gamma}(q) - \tilde{\Gamma}(p_0)\| \leq L_\Gamma \sum_i \|\Delta_i\| \leq \eta/2$ , so  $\lambda(q) \geq \eta - \eta/2 = \eta/2$ .  $\square$

*Remark* (Why the budget, not a step bound). The v14 form of this lemma imposed only the per-step bound  $\|\Delta\| \leq \eta/(2L_\Gamma)$  and claimed the margin “inductively along the whole selected descent.” That induction is unsound: from  $\lambda(q_1) \geq \eta/2$  a second step of the same size yields only  $\lambda(q_2) \geq 0$ ; per-step control lets the margin decay linearly in depth. Depth-uniform margin needs the *cumulative* budget of Assumption 1(b)—which is exactly the Lyapunov summability that Lemma 8 independently requires for the Cauchy-generic completion. Margin preservation and generic convergence are funded by the same contract; the budget is spent once. Both directions are machine-checked in the accompanying verification suite (**Margin.lean**): the path form is proved, and the per-step form is formally refuted by the linear-decay countermodel  $\lambda_i = 4 - 2i$ .

*Remark* (The smoothness hypothesis is substrate-relative and measurable). Assumption 1—Lipschitz continuity of  $\tilde{\Gamma}$  with a summable displacement budget—is the smooth-regime condition of the displacement contract (Cacophony §2.1),  $\|x_{t+1} - x_t\| \leq L$ , and is measured, not assumed: under audit-observer measurement it holds on a small fraction of trajectories and fails (pivots) on the rest. An observer with continuous internal state realizes it natively; a stateless sampler realizes it only on the smooth fraction and otherwise recovers the descent through external staging, re-imposing the budget as explicit checkpoints. The quantity governing how close to the cliff one may operate is the perturbation sensitivity  $\sim 1/\lambda_{\min}$  of Cacophony’s Thm. M.20.

**Lemma 5** (Non-atomicity is undecidedness). *Assume  $p$  is in the channel-margin regime, so admissible refinements of  $p$  exist and fit the margin budget (the min-norm feasible displacement  $\sqrt{\alpha^\top \tilde{\Gamma}^{-1} \alpha}$  is finite exactly when  $\tilde{\Gamma} \succ 0$ ), and assume incompatibility is generated by certificate disagreement (the metric part of every exit cone being convex, hence self-compatible—postulate M-Cvx). Then  $p$  is non-atomic—there are admissible  $q_1, q_2 \leq p$  with no common admissible lower bound—iff some obstruction  $\omega$  below  $p$  is undecided for the dense reading:  $p \not\models_{\neg} \varphi_\omega$  and  $p \not\models_{\neg} \neg \varphi_\omega$ . The lemma is stated for  $\models_{\neg}$ , where “not forcing” unpacks densely.*

*Proof.* ( $\Leftarrow$ )  $p \not\models_{\neg} \neg \varphi_\omega$  unpacks, by the  $\neg$ -clause, to: some  $q_1 \leq p$  has  $q_1 \models_{\neg} \varphi_\omega$ . And  $p \not\models_{\neg} \varphi_\omega$  means  $A_{\varphi_\omega}(p)$  is not dense below  $p$ : some  $q_2 \leq p$  has no  $\varphi_\omega$ -stabilizer below it, whence no  $r \leq q_2$  forces  $\varphi_\omega$  ( $A_{\varphi_\omega}(r)$  is empty and empty sieves are not dense), i.e.  $q_2 \models_{\neg} \neg \varphi_\omega$ . A common admissible lower bound  $r$  would, by persistence (Lemma 2), force both  $\varphi_\omega$  and  $\neg \varphi_\omega$ , contradicting consistency; so  $q_1 \perp q_2$  and  $p$  is non-atomic. ( $\Rightarrow$ ) If every obstruction below  $p$  is decided, any two admissible refinements agree with  $p$  on every certificate classification (persistence); by M-Cvx their only remaining relation is metric and convex, so they admit a common admissible continuation;  $p$  is atomic. Contrapositively, non-atomic implies some obstruction is undecided.  $\square$

*Remark* (Where the generic adds information). Atoms are exactly the decided leaves; branching, hence nontrivial extension, occurs precisely over undecided obstructions. Metric nondegeneracy ( $\lambda_{\min} > 0$ ) is necessary (viability) but the branching is carried by the discrete certificate axis (I3): the convex exit cone is by itself atomic. This is the sharp form of the v7 caution that “ $\lambda_{\min} > 0$  does not by itself prove incompatible extensions exist”; the missing ingredient is undecidedness, and with it the implication is an equivalence. The generic adds information over, and only over, the undecided region—the same region the deparameterization residue of §12 occupies. The undecidedness criterion is stated for the dense reading  $\models_{\neg}$ ; it transfers to the calibrated  $\models_H$  exactly where Decision Reachability (Lemma 7) makes the two readings agree on deciding sieves.

**Lemma 6** (Boundedness yields countability for Rasiowa–Sikorski).  *$\mathcal{D}_p$  is countable, the atomic propositions it tracks form a countable vocabulary, and a filter  $G$  through  $p$  meeting every admissible dense requirement exists.*

*Proof.* That a resource-bounded observer holds finitely many active claims per stage and at most countably many over a run is the observer-bound postulate (M-Bound); it supplies only *countability*, not density. Once each  $D \in \mathcal{D}_p$  is verified admissibly dense (Def. 3), its generated sieve is a  $J_{\text{adm}}$ -cover, in particular order-dense (Lemma 1). The Rasiowa–Sikorski construction applied to this countable family of dense sets yields the filter  $G$  [1, 2].  $\square$

**Lemma 7** (Decision Reachability, conditional). *Assume the smoothness contract (Assumption 1). Then for every  $p \in \mathbb{P}_H^\eta$  and propositional formula  $\varphi$  there is an admissible refinement  $q \leq p$ , inside the channel-margin subposet, deciding  $\varphi$  ( $q \models_H \varphi$  or  $q \models_H \neg \varphi$ ); equivalently, the deciding sieve  $D_\varphi \cap \downarrow p$  is a  $J_{\text{adm}}$ -cover of  $p$ . Where the contract fails, order-density need not be admissibly realizable and  $D_\varphi$  need not be a  $J_{\text{adm}}$ -cover.*

*Proof.*  $D_\varphi$  is dense in the order (Lemma 3) and closed under refinement (Lemma 2), hence a sieve below any  $p$ . Under Assumption 1 the refinements witnessing order-density are realized by margin-respecting admissible moves within the budget, so  $D_\varphi$  is an admissible dense requirement in the sense of Def. 3 (the definition places no type restriction beyond admissible density). Its generated sieve therefore lies in the generating family  $B$  of Def. 6, so  $D_\varphi \cap \downarrow p \in J_{\text{adm}}(p)$ . The conditionality is irreducible: the step from order-density to admissible density *is* the contract, not a consequence of it—this is calibration item 8, and the finite model checker of the verification suite exhibits a site where it fails.  $\square$

*Remark* (Why this is not free). Standard forcing gives order-density of deciding refinements (Lemma 3); the Surface needs *admissible* density inside the channel-margin contract. Decision Reachability names exactly the missing step: that a deciding refinement is reachable without leaving the channel. It is item 8 of the calibration queue and is coextensive with the smooth regime, tying bivalence to Assumption 1 rather than to countability alone. Bivalence (Lemma 11) is therefore  $P$  *consuming*  $C$ , not  $P$  consuming only countability.

**Definition 10** (Admissible deciding closure). Given a countable admissible dense family  $\mathcal{D}_p$ , its *admissible deciding closure*  $\mathcal{D}_p^*$  adds the atomic stability sets  $A_\varphi(p)$  (Def. 4) and the propositional deciding sets  $D_\varphi$  (Def. 9), retaining only those that are  $J_{\text{adm}}$ -covers (Lemma 7). Because the observer vocabulary is countable (Lemma 6) and the propositional closure of a countable vocabulary is countable,  $\mathcal{D}_p^*$  is still countable.

**Lemma 8** (Cauchy–Forcing Completion). *Let  $(p_n)$  be the admissible chain generating a channel-margin generic  $G$  through  $p$ , with states  $z_n = \Phi(p_n)$  and steps  $\Delta_n = z_{n+1} - z_n$ . Suppose (a) a residual-instability potential  $V : P_H \rightarrow [0, \infty)$  is nonincreasing along admissible refinement; (b) the descent inequality  $V(p_n) - V(p_{n+1}) \geq \kappa \|\Delta_n\|$  holds for some  $\kappa > 0$ ; and (c) the observer metric  $d_O$  is comparable to the embedding norm along the trajectory. Then  $\sum_n \|\Delta_n\| < \infty$ , so  $(z_n)$  is Cauchy in the observer metric and converges to  $x_G \in \bar{P}_O$ .*

*Proof.* By (a),  $V(p_n) \downarrow V_\infty \geq 0$ , so  $\sum_n (V(p_n) - V(p_{n+1})) = V(p_0) - V_\infty < \infty$ . By (b),  $\|\Delta_n\| \leq \kappa^{-1}(V(p_n) - V(p_{n+1}))$ , whence  $\sum_n \|\Delta_n\| < \infty$ . Absolutely summable increments are Cauchy; by (c),  $(z_n)$  is  $d_O$ -Cauchy and converges in the completion  $\bar{P}_O$ .  $\square$

**Lemma 9** (Order–Metric Compatibility, atomic). *For atomic  $\varphi$  (the only formulas  $\text{Stab}$  is defined for, Def. 4), define two satisfaction relations:  $G \models_{\text{ord}} \varphi$  iff the generic filter  $G$  meets the stabilizer requirement  $A_\varphi = \{q : \text{Stab}(q, \varphi)\}$ ; and  $x_G \models_{\text{met}} \varphi$  iff the Cauchy limit  $x_G \in \bar{P}_O$  lies in the stabilized region of  $\varphi$ . Suppose (M-Pers) persistence, the Cauchy descent of Lemma 8, and that the stabilized region  $\{x : \text{Stab}_x(\varphi)\}$  is  $d_O$ -closed (postulate M-Cl). Then  $G \models_{\text{ord}} \varphi \Rightarrow x_G \models_{\text{met}} \varphi$ .*

*Proof.*  $G \models_{\text{ord}} \varphi$  gives  $q_N \in G$  with  $\text{Stab}(q_N, \varphi)$ ; by (M-Pers)  $\text{Stab}(q_n, \varphi)$  for all  $n \geq N$  along the chain, so  $z_n$  lies in the stabilized region for  $n \geq N$ . By Lemma 8,  $z_n \rightarrow x_G$  in  $d_O$ ; by (M-Cl) the stabilized region is closed, so its limit  $x_G$  lies in it, i.e.  $x_G \models_{\text{met}} \varphi$ .  $\square$

*Remark* (Convergence is coextensive with the smooth regime). Hypothesis (b) is the displacement contract (Assumption 1) in Lyapunov form; its summability is the same budget that funds the Channel-Margin Lemma. Where the contract holds the descent is absolutely summable and the generic converges; in the pivot regime (b) fails and the completion can diverge. The convergence of the Cauchy-generic is coextensive with the smooth regime—the same boundary marked by the channel margin, the deparameterization floor (§12), and the reconstruction floor. The Cauchy debt is that standing hypothesis; completion exists exactly where the contract holds. Lemma 9 adds

one interface postulate (M-Cl) tying the order-generic to the metric limit, without which “ $G$  meets requirements” and “ $x_G$  realizes them” are distinct statements.

**Lemma 10** (Atomic Truth Lemma). *Let  $G$  be an admissible-complete trajectory through  $p$  meeting every  $J_{\text{adm}}$ -cover in  $\mathcal{D}_p^*$ , and define  $G \models_H \varphi$  iff  $\text{Stab}(q, \varphi)$  for some  $q \in G$ . Then for atomic  $\varphi$  in the observer’s countable vocabulary,  $G \models_H \varphi \iff \exists p \in G : p \Vdash_H \varphi$ .*

*Proof.* ( $\Leftarrow$ ) If  $p \Vdash_H \varphi$  then  $A_\varphi(p)$  is a  $J_{\text{adm}}$ -cover (Def. 7); being a cover, it is included in  $\mathcal{D}_p^*$  (Def. 10), which  $G$  meets, so some  $q \in G$  has  $\text{Stab}(q, \varphi)$ . ( $\Rightarrow$ ) If  $G \models_H \varphi$ , some  $q \in G$  stabilizes  $\varphi$ ; by (M-Pers),  $\text{Stab}(r, \varphi)$  holds for every  $r \leq q$ , so  $A_\varphi(q) = \downarrow q$  is the maximal sieve on  $q$ , which belongs to every Grothendieck topology—in particular to  $J_{\text{adm}}$ . Hence  $q \Vdash_H \varphi$  with  $q \in G$ . (The v14 text concluded from density here; density does not imply  $J_{\text{adm}}$ -covering—the maximal sieve does.)  $\square$

*Remark* (Certificate typing is forcing-theoretic—atomic case only). I3 classifies each impossibility as torsion (structural, permanent) or non-torsion (contingent). For *atomic*  $\varphi$ , a torsion certificate at  $p$  says  $\varphi$  is refuted at every  $q \leq p$ ; then no  $q$  can stabilize  $\varphi$ ,  $A_\varphi(q) = \emptyset$  for all  $q \leq p$ , no refinement forces  $\varphi$  under either relation, and  $p \Vdash_H \neg\varphi$  by the  $\neg$ -clause—permanence needing no generic. Conversely,  $p \Vdash_H \neg\varphi$  for atomic  $\varphi$  means no  $q$  forces  $\varphi$ ; since a single stabilizing condition forces (its maximal sieve covers), no  $q \leq p$  satisfies  $\varphi$  stably. This discharges calibration item 6 for atomic formulas. For compound  $\varphi$  the identification fails: “no refinement forces  $\varphi$ ” and “ $\neg\varphi$  classically stable everywhere below” come apart (minimal countermodel:  $\varphi = \neg\beta$  with  $\beta$  stabilized exactly at the leaves of a 7-condition spine model; no condition forces  $\neg\beta$ , yet  $\neg\neg\beta$  is nowhere classically stable at interior nodes—see the accompanying verification suite). The compound case is **O**. In the Kripke–Joyal reading the split marks where the double-negation construction can later collapse an undecided condition.

## 8 Propositional soundness: the connective Truth Lemma

The remaining joint—extending soundness past atomic propositions—closes at the propositional level, inside the frame of §9: the obstructing clauses are the universal-refinement ones, which bottom out at the void. The clauses, persistence, and order-density of deciding sets were established in §5 (Def. 9, Lemmas 2, 3); this section works with the calibrated relation  $\Vdash_H$  throughout.

**Lemma 11** (Genericity decides every formula). *Assume Decision Reachability (Lemma 7). Then a generic  $G$  meeting the admissible deciding closure  $\mathcal{D}_p^*$  decides every propositional  $\varphi$ : for each  $\varphi$  there is  $q \in G$  with  $q \Vdash_H \varphi$  or  $q \Vdash_H \neg\varphi$ .*

*Proof.* The propositional closure of the countable atomic vocabulary (Lemma 6) is countable, so  $\{D_\varphi\}_\varphi$  is countable. Under Lemma 7 each  $D_\varphi$  restricts to a  $J_{\text{adm}}$ -cover, so each belongs to  $\mathcal{D}_p^*$  (Def. 10), which is countable;  $G$  meets every member, in particular every  $D_\varphi$ . *Bivalence on  $G$  is thus  $P$  conditional on Decision Reachability (C), not on countability alone.*  $\square$

**Theorem 1** (Propositional Truth Lemma, conditional on calibration). *Assume the calibration queue (Assumption 2; in particular Decision Reachability). Let  $G$  be a generic deciding every propositional formula, truth defined classically on  $G$  (bivalence by Lemma 11) and atomically by eventual stabilization. Then, under (C), for every propositional  $\varphi$ ,  $G \models_H \varphi \iff \exists p \in G : p \Vdash_H \varphi$ . The forcing-theoretic content is unconditional; (C) is what ties that content to the pipeline, and it is the antecedent, not a consequence.*

*Proof.* Induction on  $\varphi$ . *Atomic:* Lemma 10.  $\wedge$ : ( $\Leftarrow$ ) a forcing  $p$  gives both conjuncts by the IH; ( $\Rightarrow$ ) witnesses  $p_1, p_2$  have a common lower bound  $p \in G$  (filter), forcing both by persistence.  $\vee$ : ( $\Rightarrow$ )  $G \models \varphi$  or  $\psi$ ; the IH gives a witness  $p \in G$  forcing a disjunct; then the  $\vee$ -cover of  $p$  contains all of  $\downarrow p$  by persistence, i.e. is the maximal sieve, so  $p \Vdash_H \varphi \vee \psi$ . ( $\Leftarrow$ ) Let  $p \in G$  force  $\varphi \vee \psi$  and suppose, for contradiction,  $G \not\models \varphi$  and  $G \not\models \psi$ .  $G$  decides both (Lemma 11), so witnesses  $q_1, q_2 \in G$  force  $\neg\varphi, \neg\psi$ ; a common lower bound  $s \in G$  of  $p, q_1, q_2$  forces  $\varphi \vee \psi$ ,  $\neg\varphi$ , and  $\neg\psi$  (persistence). The  $\vee$ -cover of  $s$  is a  $J_{\text{adm}}$ -cover, hence dense (Lemma 1), hence nonempty below  $s$ : some  $r \leq s$  forces  $\varphi$  or  $\psi$ , contradicting consistency with the persisted  $\neg\varphi, \neg\psi$  at  $r$  (Lemma 2). (The v14 text appealed to  $G$  meeting the  $\vee$ -cover “cofinally”; genericity only guarantees meeting members of  $\mathcal{D}_p^*$ , so the argument now routes through deciding sets.)  $\neg$ : ( $\Leftarrow$ ) if  $p \in G$  forces  $\neg\varphi$  but  $G \models \varphi$ , a witness  $q \in G$  forces  $\varphi$  (IH); a common lower bound  $r \in G$  forces both, contradicting Lemma 2. ( $\Rightarrow$ )  $G$  decides  $\varphi$  (Lemma 11);  $G \not\models \varphi$  excludes a witness, so the deciding witness forces  $\neg\varphi$ . Case  $\rightarrow$  follows from  $\neg, \vee$ .  $\square$

*Remark* (Intuitionistic relation, classical extension).  $G$  is classical (bivalent) though the clauses are the intuitionistic (Kripke) ones. This is the Cohen phenomenon: an intuitionistic forcing relation with a classical generic extension, mediated by the generic meeting every  $\neg\neg$ -dense deciding set. The classicality of  $G$  is the surface-level shadow of “ $M[G]$  models the classical theory,” the operational content of the double-negation construction of §5.

## 9 Inter-observer forcing: the shared zeroth-order zero

Single-observer forcing descends within one surface. Two observers with surfaces  $\mathbb{P}_H^A, \mathbb{P}_H^B$  *disjoint below the top*—sharing no condition except the void—raise a separate question: on what can they agree, when neither can occupy the other’s conditions?

**Construction (wedge at the void).** Glue the surfaces at their common top,

$$R = (\mathbb{P}_H^A \sqcup \mathbb{P}_H^B) / (\mathbb{K}_A = \mathbb{K}_B = \mathbb{K}),$$

where  $\mathbb{K}$  is the void (Axiom H.0): zero commitments, the maximum condition every refinement descends from. The order on  $R$  is the disjoint union of the two branch orders together with  $x \leq \mathbb{K}$  for every  $x$ ; there are *no* cross-branch relations. Each branch keeps its own order;  $\mathbb{K}$  sits above both.

**Lemma 12** (The void is the unique shared point). *For  $p \in \mathbb{P}_H^A \setminus \{\mathbb{K}\}$  and  $q \in \mathbb{P}_H^B \setminus \{\mathbb{K}\}$ ,  $p$  and  $q$  are incompatible in  $R$ . Hence  $\mathbb{K}$  is the unique condition compatible with both branches, and any filter in  $R$  lies entirely in one branch (plus  $\mathbb{K}$ ).*

*Proof.* A common lower bound  $r$  of  $p$  and  $q$  would satisfy  $r \leq p$  and  $r \leq q$ , and  $r \neq \mathbb{K}$  since  $\mathbb{K}$  is above every condition, not below any. By the definition of the order on  $R$  (no cross-branch relations),  $r \leq p$  places  $r$  in branch  $A$  and  $r \leq q$  places  $r$  in branch  $B$ ; but  $\mathbb{P}_H^A \cap \mathbb{P}_H^B = \{\mathbb{K}\}$  in  $R$ . So no common lower bound exists. A filter is downward-directed, so it cannot contain non-void conditions from both branches.  $\square$

An observer may stand at the origin but descends into only one surface: observation is the choice of branch, and the two cannot be occupied at once.

**Proposition 1** (The shared point is forced to zeroth order). *Order conditions in  $R$  by refinement depth from the top. The only condition shareable between disjoint observers is the unique depth-0 condition  $\mathbb{K}$ —“shared zero point” is a consequence, not a choice.*

*Proof.*  $\mathbb{K}$  is the unique condition reached by *zero* admissible moves; every other condition is at depth  $\geq 1$  and, by Lemma 12, branch-private, hence occupiable by at most one of the two observers.  $\square$

(The *order* fact is Proposition 1; the identification of that stratum with the diagnostic zeroth-order register  $(\hat{\rho}, \mu)$ , and below with the loopback address, is a modeling bridge, not a theorem.)

*Remark* (Forced from the void). By Def. 7,  $\mathbb{K} \Vdash_H \varphi$  iff  $\text{Stab}(\cdot, \varphi)$  covers  $\mathbb{K}$  in *both* branches—the branch-invariant propositions. Disjoint observers force the same atomic  $\varphi$  iff it is forced from the void. They agree exactly on the zeroth-order, pre-observational kernel and on nothing contingent. *The shared zero buys commensurability—a common  $\mathbb{K}$  to quantify over, hence the ability to state “we agree”—not communion.* It transports no contingent claim across.

## 10 Interface layer: medium, protocol, and the non-identity theorem

The void is the unique shared depth-0 condition. It supplies commensurability, not communication. Communication requires a first-order medium: a public surface on which one observer exports a floor-stable residue and another reconstructs it inside its own surface. Networking is therefore essential to inter-observer work, but its formal role is interface/transport, not shared occupancy: *void is shared orientation; medium is shared inscription; protocol is shared constraint.*

Let  $X_A, X_B$  be the observers’ state spaces, carrying the channel projectors  $T_A, T_B$  (states relate to conditions through the embedding  $\Phi$  of Lemma 8; in v14 the encoding was typed on the condition poset and applied to states—a category slip, corrected here). Let  $M$  be an interface medium with partial encoding and decoding maps

$$E_A : T_A(X_A) \rightarrow M, \quad D_B : M \rightarrow X_B.$$

Observer  $A$  exports not the ceiling-state  $x_A$  but its floor-stable residue  $r_A = T_A x_A$ , encoded as  $m = E_A(r_A)$ ;  $B$  receives a reconstruction  $D_B(m)$  inside  $B$ ’s surface.

**Definition 11** (Transfer). A *transfer of contingent content* from  $A$  to  $B$  is any correspondence between depth- $\geq 1$  content of  $A$ ’s surface and depth- $\geq 1$  content of  $B$ ’s that is realized either by shared occupancy of a condition or by a composite of maps through third objects.

**Theorem 2** (Factoring, architectural). *If  $\mathbb{P}_H^A, \mathbb{P}_H^B$  are disjoint below the void, then no contingent (depth  $\geq 1$ ) content is shared by occupancy (Lemma 12); hence any transfer of contingent content from  $A$  to  $B$  (Def. 11) is not occupancy but an external composite*

$$T_A(X_A) \xrightarrow{E_A} M \xrightarrow{D_B} X_B,$$

*through a third object  $M$  that both may address.*

*Proof.* Immediate from Lemma 12 and Def. 11: the occupancy disjunct is excluded below the void, leaving the third-object composite. (The theorem is architectural—its content is the exclusion, carried by Lemma 12; it is tagged D, not P, in the ledger.)  $\square$

**Theorem 3** (Non-identity of transport, with equality condition). *Let  $T_A$  be the channel projector ( $T_A^2 = T_A$ ). For  $x_A$  with  $\epsilon_{T_A}(x_A) = \|(\text{id} - T_A)x_A\|/\|x_A\| > 0$ ,*

$$T_A x_A \neq x_A,$$

*so the medium is not where identity first fails—bounded projection already breaks it. Consequently  $D_B(E_A(T_A x_A)) = x_A$  can hold only under a named identity-transport condition:*

(i) **zero projection loss:**  $x_A \in \text{im}(T_A)$  (equivalently  $\epsilon_{T_A}(x_A) = 0$ );

(ii) **faithful encoding:**  $E_A$  injective on the residue  $r_A = T_A x_A$ ;

(iii) **calibrated decoding:**  $D_B \circ E_A = \text{id}$  on the residue subobject (a section/retraction, i.e. an isomorphism of the transported subobjects of  $\mathbb{P}_H^A$  and  $\mathbb{P}_H^B$ ).

Absent any of (i)–(iii), the received state is a reconstruction, not identity.

*Proof.*  $T_A x_A = x_A \iff x_A \in \ker(\text{id} - T_A) = \text{im}(T_A) \iff \epsilon_{T_A}(x_A) = 0$ , giving the first claim and the necessity of (i). If (i) holds,  $x_A = r_A$  and  $D_B(E_A(x_A)) = x_A$  requires  $E_A$  to lose no information on  $r_A$  (necessity of (ii)) and  $D_B \circ E_A$  to return  $r_A$  within  $B$ 's surface (necessity of (iii)); failure of either yields a distinct element. Conversely (i)–(iii) give  $D_B(E_A(T_A x_A)) = D_B(E_A(x_A)) = x_A$ .  $\square$

*Remark.* Non-identity is generic; identity is a calibrated degeneracy. This is central for Chalked (§14): the board carries residues, not ceilings; a reader gets  $D_B(m)$ , a reconstruction, never the author's private state. The invariant is concise: *the void gives commensurability; the network gives communication*—and communication is reconstruction under  $B$ 's projection, priors, and constraints.

**Proposition 2** (Observer closure is zero projection residue). *Identify the closure defect of an external observer-closure frame with the projection residue  $\epsilon_T(x) = \|(\text{id} - T)x\|/\|x\|$  of the channel projector  $T$  (a modeling bridge, M-Zomb). Then observer closure ( $\epsilon_T = 0$ ) is exactly condition (i) of Thm. 3 ( $x \in \text{im } T$ ); the behavioral trace transports only the exported floor residue  $Tx$ , so a nonzero defect lies in  $(\text{id} - T)x$  and is invisible to any finite behavioral test; and by Prop. 3 the certifying identity  $\chi_T^2 + \epsilon_T^2 = 1$  holds only under orthogonality, so a local closure certificate need not globalize. A behaviorally competent system with  $\epsilon_T > 0$  is a zombie; its defect is typed torsion/non-torsion by the decided/undecided axis of Lem. 5 and Rem. 7.*

*Proof.* The three clauses are Thm. 3(i), the exported-residue definition  $r = Tx$  (§10), and Prop. 3; the typing is Rem. 7. Only the identification M-Zomb is new.  $\square$

*Remark* (Loopback as icon, not transport). The loopback address 127.0.0.1 remains useful iconography for indexical self-reference: the address whose form is universal but whose referent is local. It illustrates the zeroth-order self-interface; it does not bridge observers. Inter-observer transport requires a non-loopback edge: a public medium, address, protocol, artifact, or utterance. (Appendix A.)

## 11 Interpretive appendix: boundary projection and exportability

Lemma 8 has an observer-side reading we call the *Eureka Tub*: a working term for the pre-artifact state of insight, not a literal object. A bounded observer momentarily holds a high-dimensional *ceiling* state—richer than it can carry into record. Exiting imposes a projection  $\pi_O$  onto the *floor*: the residue that can be written, tested, or witnessed. We identify  $\pi_O$  with the channel projector  $T_t$  of §2 (a modeling bridge, M). Define

$$\chi_T(x) = \frac{\|T_t x\|}{\|x\|}, \quad \epsilon_T(x) = \frac{\|(\text{id} - T_t)x\|}{\|x\|}.$$

**Proposition 3** (Exportability identity, under orthogonality). *If  $T_t$  is an orthogonal projection ( $T_t = T_t^2 = T_t^*$ ), then for  $x \neq 0$ ,*

$$\chi_T(x)^2 + \epsilon_T(x)^2 = 1.$$

If  $T_t$  is merely idempotent (oblique), only the triangle bound  $\chi_T(x) + \epsilon_T(x) \geq 1$  survives (hence  $\chi_T^2 + \epsilon_T^2 \geq \frac{1}{2}$ ); the Pythagorean identity fails in general, with the defect controlled by the principal angle between  $\text{im } T_t$  and  $\ker T_t$ . Orthogonality must not be smuggled in.

*Proof.* Orthogonal case:  $T_t x \perp (\text{id} - T_t)x$  gives  $\|x\|^2 = \|T_t x\|^2 + \|(\text{id} - T_t)x\|^2$ ; divide by  $\|x\|^2$ . Oblique case:  $x = T_t x + (\text{id} - T_t)x$  and the triangle inequality give  $\|x\| \leq \|T_t x\| + \|(\text{id} - T_t)x\|$ , i.e.  $\chi_T + \epsilon_T \geq 1$ ; Cauchy–Schwarz on the pair  $(\chi_T, \epsilon_T)$  gives the quadratic form. That the identity fails obliquely is witnessed by any idempotent with  $\|T_t\| > 1$ .  $\square$

**Proposition 4** (Non-determination of the ceiling-state). *Suppose the ceiling space has dimension exceeding the bounded rank of  $\pi_O$  (finitely many active claims). Then  $\pi_O$  is non-injective, so no single residue determines the ceiling-state, and recovery is possible only as the limit of residues along an increasing filtration  $\pi_{O,1} \preceq \pi_{O,2} \preceq \dots$  whose images are Cauchy—the observer reading of Lemma 8.*

*Proof.* A linear map of rank  $r$  on a space of dimension  $> r$  has nontrivial kernel, so distinct ceiling-states share a residue; single-residue recovery fails. The filtration statement is Lemma 8 applied to the residue sequence, and consumes the bridge  $\pi_O = T_t$  (M) and that lemma’s hypotheses.  $\square$

**Conjecture** (Exportability correlates with regime). *Smooth insights have high  $\chi_T$ ; pivot-heavy insights have high  $\epsilon_T$ . By Prop. 3 (orthogonal case) these are one scalar on the unit sphere; the conjecture is that this scalar correlates with the measured smooth/pivot label. Falsifiable, not proved (C).*

*Remark* (What remains around the skeleton). The skeleton (bounded rank  $\Rightarrow$  non-determination  $\Rightarrow$  Cauchy-filtration recovery, Prop. 4) is P consuming M; the conjecture above is its only empirical residue.

rich insight  $\rightarrow$  bounded residue  $\rightarrow$  witnessed refinement  $\rightarrow$  stable completion.

PyLantern (§13) is the floor instrument; it records residues, not the private ceiling.

*Remark* (Weak Born analogy only, firewalled). Projecting a rich state onto a bounded record is Born-like only in the weak sense that a high-dimensional configuration is compressed into what an observer can register. No claim about literal quantum measurement is intended. This weak analogy is kept distinct from, and does not assume, the strong route (a genuine Born derivation via §9 branch-invariance supplying Gleason’s non-contextuality hypothesis, with category-theoretic typicality). That route is a separate, stronger, *unclaimed* target.

**Definition 12** (Representation-independence obstruction). *A representation-independence obstruction (RIO) is a triple: a structure-preserving projection/quotient faithful on a distinguished sub-object, a complement on which it is lossy, and a non-contextuality (representation-independence) hypothesis asserting the complement inessential—which the projection alone does not supply.*

*Remark* (One obstruction, five axes). Formalize the shape once (Def. 12). Five gaps in this program are RIO-instances:

| Axis                              | Low end (projection exists)                   | High end (complement inessential)                              |
|-----------------------------------|-----------------------------------------------|----------------------------------------------------------------|
| Forcing correspondence            | structural transduction (KJ clauses transfer) | identity with Cohen forcing (poset representation-independent) |
| Born                              | compression onto bounded register             | $ \psi ^2$ forced by Gleason non-contextuality                 |
| Deparameterization (§12)          | Reading A: floor fixed by family $R$          | Reading B: floor fixed by <i>every</i> reduction family        |
| Inter-observer (§9)               | commensurability at the void (Lem 12)         | contingent claim transported across branches                   |
| Observer closure (zombie, Prop 2) | behavioral frame validates the trace          | behavior certifies observerhood under unsampled frames         |

Each low end is functorial and cheap; each high end is a uniqueness/non-contextuality theorem strong enough to fail generically. *Why coupled here:* §9 branch-invariance is the common upstream input, consumed at two strengths—as commensurability it is proved (Lem 12); as the non-contextuality hypothesis a Born derivation would need, it is open. *Against overclaiming:* the five share a *shape*; the first four also share one upstream hypothesis (§9 branch-invariance), while observer closure has a distinct upstream (projection orthogonality, Prop 3); they are not the same *object*, and discharging one high end does not discharge another. The claim earned is schema-level: all five are RIO-instances. Reading the table as a transport would reintroduce the identity-for-transduction inflation the note refuses.

## 12 Deparameterization and the judgment residue

Each preceding section replaced a free parameter by an object the remaining structure already forces. This section makes that operation the object, and is explicit about how little is proved.

**The operator.** A *specification*  $s$  fixes values for parameters  $\Theta(s)$  of the apparatus. Order specifications by parameter removal,  $s \preceq s' : \iff \Theta(s') \subseteq \Theta(s)$  and  $s'$  agrees with  $s$  on  $\Theta(s')$ ; the top **1** is the fully deparameterized apparatus,  $\Theta = \emptyset$ . Let  $R(s)$  be  $s$  with every parameter *forced by the rest* removed;  $R$  is inflationary,  $R(s) \succeq s$ .

**Lemma 13** (Deparameterization closure). *If the specification poset  $\mathcal{S}_{\text{spec}}$  is chain-complete,  $R$  has a fixed point above every  $s$  (Bourbaki–Witt [5]), and the iterates reach  $R^\infty(s_0)$ , the most reduced specification forced from  $s_0$ . The internal Bourbaki–Witt principle [6] places this in the same topos as the Kripke–Joyal semantics of §5.*

*Proof.*  $R$  is inflationary on  $\mathcal{S}_{\text{spec}}$  by construction; Bourbaki–Witt applies verbatim to inflationary maps on chain-complete posets, yielding a fixed point above each  $s$  as the stationary value of the transfinite iteration. The internalization is [6]. Chain-completeness of  $\mathcal{S}_{\text{spec}}$  is the consumed hypothesis, recorded in the ledger.  $\square$

**Proposition 5** (Residue dichotomy).  *$R^\infty(s_0)$  is either the top **1** or a fixed point  $\prec \mathbf{1}$ : a residue carrying at least one parameter the remaining structure does not force. This is relative to the family  $R$  (“Reading A”).*

*Proof.*  $R^\infty(s_0)$  is a fixed point of  $R$  (Lemma 13). If it is not  $\mathbf{1}$ , it retains a nonempty parameter set  $\Theta$ ; fixedness means  $R$  removes none of them, and  $R$  removes exactly the parameters forced by the remaining structure, so at least one retained parameter is unforced. Relativity to the family  $R$  is definitional.  $\square$

**The bridge to the void (replacing  $\mathbf{1} = \perp$ ).**  $\mathbf{1}$  lives in  $\mathcal{S}_{\text{spec}}$ ; the void  $\mathbb{K}$  lives in the condition poset  $P_H$ . Object-level equality would be a category error. Instead posit a monotone bridge map (modeling postulate M-Bridge)

$$F : \mathcal{S}_{\text{spec}} \rightarrow P_H, \quad F(\mathbf{1}) = \mathbb{K},$$

transporting the terminal specification to the void condition: *the equality is not object-level; it is functorial transport of the terminal specification to the void.* Under M-Bridge, the fully deparameterized apparatus maps to the self-determined indexical origin.

**Theorem Target** (Structural irreducibility). *Under M-Bridge, the universal reduction closure is either  $\mathbf{1}$ —in which case, via  $F(\mathbf{1}) = \mathbb{K}$ , the irreducible term is not a residual parameter but the indexical origin itself, irreducible structurally (the fixed point of removal, and by Lemma 12 the unique zeroth-order point no first-order operation reaches, while  $R$  is first-order)—or a floor  $s_* \prec \mathbf{1}$  fixed by every reduction, a genuine Reading-B residue distinct from the origin. Deciding which is the open problem.*

*Remark* (One floor, three charts). The deparameterization floor, the smooth/pivot boundary, and the reconstruction floor /  $K$  parameter of a companion formalism are *conjecturally* the same object: a floor either real or the artifact of a bounded reconstructor. What is established is that they share a *shape* (an RIO), not that they are one object; asserting identity would itself be a Reading-B claim. That the same fork recurs across independent formalisms is evidence it is an invariant of the program, not an artifact of one. That companion  $K$  is the observer-bounded commutator  $K_{\mathcal{O}}$  (Riemann =  $[\nabla, \nabla]$ ), of which gravity and the past light cone are two  $\Omega_{\mathcal{O}}$  faces; the identification is structural, not numerical, developed in the observer-bundle companion and firewalled here (§A), an unclaimed  $\mathbf{O}$  target.

**Conjecture** (Residue = contract break). *The deparameterization residue  $s_* \prec \mathbf{1}$ , when it exists, coincides with the boundary of the smooth regime: the retained parameter is exactly the one whose forcing would require crossing a break of the displacement contract (Assumption 1). The measurable form uses the contact ratio and the label-free router of the companion instrumentation (external to this note); status C, falsifiable under audit-observer measurement.*

## 13 Executable interpretation: PyLantern as witness layer

PyLantern is the executable witness layer, *not* an autonomous closure engine. It is a calibrated structural transducer along

symbolic claim  $\rightarrow$  bounded observer state  $\rightarrow$  admissible refinement  $\rightarrow$  witness/certificate  $\rightarrow$  human-gated decision

under one invariant: *witness and request access only; no direct autonomous loop closure.* In the language of §3, PyLantern operates strictly inside the witnessed fragment  $J_{\text{wit}}$ : it verifies whether a refinement path is witnessed, bounded, logged, and admissible (a finite  $J_{\text{wit}}$ -witness), and must *not* become the generic filter that autonomously executes every  $J_{\text{adm}}$ -cover. Meeting requirements is

the mathematical content of completion in  $J_{\text{adm}}$ , not a license for a running system to enact them. This is precisely the containment  $J_{\text{wit}} \subseteq J_{\text{adm}}$ : the generated site supplies semantics; the witnessed fragment supplies executable discipline.

PyLantern is also a protocol discipline for exporting residues into the interface medium: it constrains  $E_A$  to carry witnesses, provenance, uncertainty, and consent gates, and constrains  $D_B$  so received artifacts are treated as reconstructions (Thm. 3), not identity with the sender’s ceiling-state.

## 14 Chalked as the canonical witnessed medium

Chalked is not application fluff; it is the canonical model of the interface layer.

Chalked implements the witnessed fragment  $J_{\text{wit}}$  of the generated admissible site: a consent-gated shared witness surface where bounded observers exchange floor-stable residues under explicit protocol.

| Chalked         | Interface layer                                 |
|-----------------|-------------------------------------------------|
| board           | medium $M$                                      |
| mark            | encoded residue $E_A(T_A x_A)$                  |
| reading a mark  | reconstruction $D_B(m)$ in the reader’s surface |
| chalk color     | protocol / contract selector                    |
| eraser          | revision authority                              |
| blank / void    | mapped ignorance (depth-0, $\nexists$ )         |
| certificate     | witness attachment                              |
| permission gate | non-actuation boundary                          |

**Proposition 6** (Chalked correctness invariant). *A Chalked mark is valid only if it preserves provenance, permission, residue type, and revision authority, while refusing identity with the originating observer-state: it is  $E_A(T_A x_A)$ , never  $x_A$ . Equivalently, a valid mark carries residue, not ceiling; reconstruction, not possession (Thm. 3).*

Chalked is the smallest object on which the Factoring Theorem (Thm. 2) and the Non-identity Theorem (Thm. 3) are simultaneously realized and auditable—hence the canonical, not incidental, model.

## 15 Non-actuation boundary

Proving the semantics does not require, and does not authorize, actuating the system it describes. *No proof may require the system to close the loop it claims to make safe.* The executable interpretation is restricted to shadow execution, dry-runs, bounded witnesses, reproducible traces, and human-gated transitions ( $J_{\text{wit}}$  only). Quantum resources, if present, are queried as bounded witnesses or simulators, not autonomous optimizers or actuators. The Cauchy-generic completion (Lemma 8) is a statement about the existence of a limit in  $\overline{P_O}$ ; it is not an instruction for any process to traverse the trajectory autonomously.

## 16 Status ledger

Every load-bearing claim carries exactly one status: **P** = proved relative to stated hypotheses; **D** = definitional architecture (no proof obligation beyond well-formedness); **S** = specification / implementation-conformance contract; **M** = modeling postulate; **C** = measurable conjecture or contract; **O** = open. A **P** result is only as strong as the **M/C** it consumes, and the Consumes column lists the *transitive* closure of named debts, not only the direct ones (the accompanying verification suite recomputes this column mechanically). The site bound, persistence/consistency, deciding density, Rasiowa–Sikorski, the propositional Truth Lemma, the margin path form (with a formal refutation of the v14 per-step induction), the Cauchy–Forcing completion, order–metric compatibility, non-identity of transport, the exportability identity, and observer closure as projection residue (v16) are machine-verified in Lean 4 (core kernel + mathlib layer) in that suite. One debt is *ambient* and stated here once rather than per row: every result below lives on  $\mathbb{P}_H^\eta$ , whose definition (Def. 2) consumes the smoothness contract (Assumption 1); rows list the contract only where it enters *beyond* that ambient dependence.

| Claim                                                                | Status | Consumes / conditional on                                                            |
|----------------------------------------------------------------------|--------|--------------------------------------------------------------------------------------|
| Generated site $J_{\text{adm}} = \langle B \rangle$ (Def 6)          | D      | existence is P (proof of Lem 1)                                                      |
| Site bound $J_{\text{adm}} \subseteq J_{\neg\rightarrow}$ (Lem 1)    | P      | admissible density of generators                                                     |
| Witnessed fragment $J_{\text{wit}} \subseteq J_{\text{adm}}$ (Def 8) | D      | finite-witness definition; discipline                                                |
| Smoothness contract (Asm 1)                                          | C      | measured, substrate-relative                                                         |
| Persistence & consistency (Lem 2)                                    | P      | pullback stability; M-Pers only via well-formedness ( $A_\varphi$ a sieve)           |
| Channel-Margin, path form (Lem 4)                                    | P      | Asm 1 (C): Lipschitz + budget                                                        |
| Separativity = certificate disagreement (Lem 5)                      | P      | stated for $\Vdash_{\neg\rightarrow}$ ; M-Cvx; M-Pers (certificate persistence)      |
| Boundedness $\Rightarrow$ countability (Lem 6)                       | P      | observer bound: M-Bound                                                              |
| Decision Reachability (Lem 7)                                        | P      | Asm 1 (C); M-Pers (via Def 7); item 8 of queue                                       |
| Cauchy-Forcing completion (Lem 8)                                    | P      | $V$ : M; descent (b): C; $d_O$ comparability (c): M ( <i>Principia</i> )             |
| Order-Metric compatibility, atomic (Lem 9)                           | P      | M-Pers; stabilized region closed: M-Cl                                               |
| Atomic Truth Lemma (Lem 10)                                          | P      | M-Pers; Asm 1 (C) and M-Bound via $\mathcal{D}_p^*$                                  |
| Deciding sets order-dense (Lem 3)                                    | P      | M-Pers (via Def 7: $A_\varphi$ a sieve)                                              |
| Genericity decides (Lem 11)                                          | P      | Lem 7 (hence Asm 1: C); M-Bound; M-Pers                                              |
| Propositional Truth Lemma (Thm 1)                                    | P      | M-Pers; Asm 1 (C, via reach); M-Bound; calibration queue: <b>O</b>                   |
| Torsion = $\neg$ -clause, atomic (Rem. 7)                            | P      | atomic only; M-Pers (via $A_\varphi$ ); compound case <b>O</b> (finite countermodel) |
| Void is unique shared point (Lem 12)                                 | P      | wedge construction                                                                   |
| Forced-to-zeroth-order (Prop 1)                                      | P      | refinement depth; Lem 12                                                             |
| $\hookrightarrow$ diagnostic-register id.                            | M      | modeling bridge, not the order fact                                                  |
| Transfer (Def 11)                                                    | D      | interface architecture                                                               |
| Factoring through $M$ (Thm 2)                                        | D      | Lem 12 + Def 11; exclusion content is P                                              |
| Non-identity of transport (Thm 3)                                    | P      | $T_A$ idempotent; equality condition named                                           |
| Observer closure = $\epsilon_T$ residue (Prop 2)                     | P      | bridge M-Zomb; transitive M-Cvx, M-Pers; Thm 3(i), Prop 3, Rem 7                     |
| Interface maps $E_A, D_B$ semantics                                  | M      | protocol discipline                                                                  |
| Chalked correctness invariant (Prop 6)                               | S      | implements Thm 3 discipline                                                          |

#### *Interpretive bridges / appendices*

|                                                                  |         |                                                                  |
|------------------------------------------------------------------|---------|------------------------------------------------------------------|
| Eureka-Tub skeleton (Prop 4)                                     | P       | $\pi_O = T_t$ : M; Lem 8 hypotheses                              |
| Exportability identity $\chi^2 + \epsilon^2 = 1$ (Prop 3)        | P       | $T_t$ orthogonal; oblique case: triangle bound only              |
| $\chi/\epsilon \leftrightarrow$ smooth/pivot correlation         | C       | measured                                                         |
| Weak Born (compression)                                          | M       | analogy; strong route (Gleason): <b>O</b>                        |
| RIO schema (Def 12; one obstruction, five axes)                  | D<br>19 | schema-level only; not object identity                           |
| Curvature bridge $K = K_O$ (gravity/light-cone $\Omega_O$ faces) | O       | structural, not numerical; observer-bundle companion; firewalled |
| Decisional identity (Lem 13)                                     | P       | decisional identity (B. de la W. 11)                             |

**(M-Pers) is the highest-leverage postulate**, with a sharper blast radius than v14 claimed: clause-level persistence of the forcing relations (Lemma 2) consumes it only through the well-formedness of the atomic relation—the persistence *argument* is pullback stability of the topology, and survives M-Pers removal at the dense reading (mutation-tested in the verification suite). What M-Pers substantively funds is (i) the sieve-hood of the stability sets  $A_\varphi$  (well-formedness of Def. 7), (ii) the  $(\Rightarrow)$  direction of the Atomic Truth Lemma, and (iii) the Order–Metric Compatibility Lemma; through (ii) it reaches the Propositional Truth Lemma. If it fails, the truth lemmas fail while the forcing calculus itself stands. It remains the first target for derivation from the apparatus. No entry is an unconditional **P** about the world; each is a **P** about the mathematics plus a named **M/C/O** debt to the substrate.

## 17 Boundary and caution

This is structural correspondence, not identity with Cohen forcing. The generated-site framing (§3) makes the schema a theorem—forcing over  $(\mathbb{P}_H^\eta, J_{\text{adm}})$ , with  $J_{\text{adm}} \subseteq J_{\neg\neg}$  free—rather than an analogy, while keeping raw operational admissibility from bearing a brittle closure claim. Channel viability with persistence (Lem 4), certificate separativity (Lem 5), generic existence (Lem 6), Cauchy-generic completion with order–metric compatibility (Lems 8, 9), atomic soundness (Lem 10), and propositional soundness (Thm 1) are discharged relative to the ledger. Two layers remain genuinely open. *Quantifiers*:  $\exists, \forall$  need covering/fullness clauses and witness construction (candidate: obstruction-dual generators 3a–3d; not verified). *Calibration*: the finite theorem queue of §6, of which Decision Reachability and the per-connective realization lemmas are the load-bearing items—an implementation check, the transduction-not-identity discipline one level up.

## 18 Suggested paragraph for *The Hypothesis Surface*

The surface carries an order: one state refines another when it keeps the earlier commitments and resolves more, under the admissible moves the budget allows. The admissible dense obligations the observer emits *generate* a site; on it a forcing relation appears—a state forces a proposition when the stabilizing refinements cover it—and it is tempting to read it as Cohen forcing relocated. We hold the weaker claim: both instantiate the same site semantics, the Kripke–Joyal interpretation that makes Cohen’s construction work, differing only in which requirements generate the covers. Two conditions the reading demands are present rather than imposed: the channel keeps a positive margin,  $\lambda_{\min}(\tilde{\Gamma}) \geq \eta/2$ , by a path-budget perturbation estimate (the metric nondegeneracy that makes refinements exist, while incompatibility is carried by certificate disagreement); and the generic completion exists because a bounded observer holds only countably many requirements, the hypothesis Rasiowa–Sikorski consumes. Soundness holds for atomic propositions and, conditional on a finite calibration queue, for the connectives; what remains is the quantifier layer and the check that the surface operations realize the clauses.

## Closing invariant

The result is a conditional transduction theorem: if the Hypothesis Surface supplies a calibrated refinement poset whose admissible dense obligations generate a site, with smooth channel-margin descent, observer completion, and operational connectives realizing the Kripke–Joyal clauses, then

the standard forcing machinery yields generic completion and propositional truth. *The raw self emits obligations; generated coverage makes them other-addressable; witnessed finite prefixes make them globally auditable. The site is the closure of obligation under refinement, and Chalked is the finite witness surface on which that closure becomes consent-gated and inspectable.* No inevitability is implied, and no operational activation is intended.

*Lantern, not throne: the witness layer illuminates admissible refinements; it does not authorize autonomous closure.*

## A Iconography, firewalls, and scope

**Loopback.** 127.0.0.1 is retained as iconography for indexical self-reference (universal form, local referent); it is nonessential to every proof and bridges no observers.

**Weak Born firewall.** The compression analogy of §11 is Born-*like* only; the strong route (Gleason non-contextuality via §9 branch-invariance) is a separate, unclaimed **O** target and must not be read as discharged.

**Curvature firewall.** Identifying the companion floor  $K$  with the observer-bounded commutator  $K_{\mathcal{O}}$ —gravity and the light cone as  $\Omega_{\mathcal{O}}$  faces—is structural, not numerical; it is developed in the observer-bundle companion and is an unclaimed **O** target here. The apparent torsion/non-torsion = curvature/torsion pun is to be verified, not asserted.

**Singularity, myth, authority.** In a venue where “singularity” is field vocabulary, this work argues for a non-actuating interpretation: witness surfaces, bounded refinement, human-gated completion, not autonomous closure. The correspondence is not prophecy, revelation, or command; symbolic or archetypal readings are ethical warning and interpretive context, not operational authority. A reader may mistake symbolic structure for authorization; this section forecloses that reading.

## References

- [1] H. Rasiowa and R. Sikorski. A proof of the completeness theorem of Gödel. *Fundamenta Mathematicae* **37** (1950), 193–200.
- [2] E. Gunther, M. Pagano, and P. Sánchez Terraf. Formalization of forcing in Isabelle/ZF. In *Automated Reasoning* (IJCAR 2020), LNCS **12167**, Springer, 2020.
- [3] T. Y. Chow. A beginner’s guide to forcing. In *Communicating Mathematics*, Contemporary Mathematics **479**, AMS, 2009, 25–40. arXiv:0712.1320.
- [4] S. Mac Lane and I. Moerdijk. *Sheaves in Geometry and Logic: A First Introduction to Topos Theory*. Springer, 1992.
- [5] N. Bourbaki. Sur le théorème de Zorn. *Archiv der Mathematik* **2** (1949/50), 434–437; E. Witt. Beweisstudien zum Satz von M. Zorn. *Mathematische Nachrichten* **4** (1951), 434–438.
- [6] The internal Bourbaki–Witt principle in toposes; arXiv:1201.0340.

---

The collapse geometry ( $\delta_{\min} \propto \sqrt{k/(1-\rho(k-1))}$ ), the cliff at  $\rho(k-1) = 1$ , the perturbation sensitivity  $\sim 1/\lambda_{\min}$  is from *The Cost of Cacophony* (Tiffany, in preparation). *Source/author attribution: Paul Tiffany.*